

Denis Ostroushko - HW3

Introduction

Imputation and Variable Processing

For the purpose of this assignment we retain the same imputation schemes we used in the previous two assignments. We will use imputation with the median of observed values and replace missing values with modes for categorical predictors.

Since we will need to create two models that have all possible confounders in the data set, we need to be careful with variable inclusion. We drop variable `Hisp` because it is highly correlated with other variables that contain race and ethnic information. We also drop `Drug.Add` due to the issues with its imputation. When imputed with the most common level “No”, which indicated no drug use, this variable has one unique level. Such zero variance predictors can cause problems with fitting models, so we will avoid using it in our analyses.

Variables `BMI`, `BL.Cig.Day`, `BL.Drks.Day`, `N.living.kids` are imputed with medians like `N.prev.preg`, `Birthweight` in the previous assignments.

`Use.Alc` is imputed with a mode like `Race_ethnicity`, `Use.Tob`

Problem 1

Instead of writing equations for each model, I provide formatted code chunks that contain model statements.

Logistic regression - A

```
logistic_regression_a <-  
  glm(  
    I(data$Group == "T") %>% as.numeric() ~  
  
    Race_ethnicity + Public.Asstce +  
    Use.Tob + Live.PTB +  
  
    N.prev.preg + BL.GE + BL..BOP + BL..PD.4 + BL..CAL.3,  
  
    data = data,  
    family = "binomial"  
  )
```

Logistic regression - B

For these non-linear terms, I selected polynomial degrees that minimize AIC, to some reasonable degree.

```
logistic_regression_b <-  
  glm(  
    I(data$Group == "T") %>% as.numeric() ~  
  
    Race_ethnicity + Public.Asstce +  
    Use.Tob + Live.PTB +  
  
    poly(N.prev.preg,2) +  
    poly(BL.GE,5) +  
    poly(BL..BOP,5) +  
    poly(BL..PD.4,5) +  
    poly(BL..CAL.3, 2),  
  
    data = data,  
    family = "binomial"  
  )
```

Logistic Regression - C

I decided to use forward selection. After a few trials I determined that forward variable selection includes less variables into the final output, while providing similar AIC. Also, when I compared SMD balance, forward stepping variable selection resulted in smaller SMD of the balanced covariates sets.

```
logistic_regression_c_full <-  
  glm(  
    I(data$Group == "T") %>% as.numeric() ~  
  
    (Race_ethnicity + Public.Asstce +  
     Use.Tob + Live.PTB +  
  
     N.prev.preg + BL.GE + BL..BOP + BL..PD.4 + BL..CAL.3)^2,  
  
    data = data,  
    family = "binomial"  
  )  
  
logistic_regression_c_lower <-  
  glm(I(data$Group == "T") %>% as.numeric() ~ 1,  
      data, family = "binomial")  
  
logistic_regression_c <-  
  MASS::stepAIC(logistic_regression_c_lower,  
                direction = "forward",  
                trace = 0,  
                scope = list(upper = logistic_regression_c_full,  
                             lower = logistic_regression_c_lower)  
                )
```

Logistic Regression - D

```
# remove variables that are other outcomes, or IDs  
data2 =  
  data %>%  
    select(- PID, - Birth.outcome, -  
          GA.at.outcome, -Preg.ended...37.wk,
```

```

      -Birthweight) %>%
mutate(Race_ethnicity = as.factor(Race_ethnicity))

logistic_regression_d_full <-
  glm(
    I(data$Group == "T") %>% as.numeric() ~ .,
    data = data2,
    family = "binomial"
  )

logistic_regression_d <-
  MASS::stepAIC(
    logistic_regression_c_lower,
    direction = "forward",
    trace = 0,
    scope = list(upper = logistic_regression_d_full,
                  lower = logistic_regression_c_lower)
  )

```

Flexible Regression - Random Forest

```

rf = randomForest(
  Group ~ .,
  data = data2 ,
  ntree = 1000
)

```

Having obtained five sets of propensity scores, we can create five sets of weighted SMDs for all covariates, i.e. potential confounders. We contrast five sets of SMDs and weighted SMD on Figure 1.

A black vertical line is set at 0.1, as we consider any SMD beyond 0.1 to be a sign of possible imbalance.

A blue horizontal line separates covariates we explicitly called in some models from the rest of confounding values. We expect that covariates not explicitly used in confounding control will exhibit higher levels of imbalance. It appears that this is true. It is especially true for logistic regression models A and B, which used only some predictors. To reiterate, all those predictors are listed below the vertical blue line.

Lastly, due to some fitted probabilities being equal to 1 or 0, I replace these extreme values with numbers that are close to them. For 0, the replacement is $0.1 * 10^{-15}$, while for 1 the replacement is $1 - 0.1 * 10^{-15}$.




Figure 1: Comparion of confounding adjustment methods and their impact on SMDs

Part B

In order to estimate Treatment Effect using a doubly robust, I will propensity scores from logistic regression specified in part A. Additionally, I specify two outcome models, one for each outcome of interest. Definition of two models are given in the code chunks below:

```
pregnancy_model <-  
  glm(  
    `Preg.ended...37.wk` ~  
    Group + Race_ethnicity + Public.Asstce + Use.Tob + N.prev.preg +  
    Live.PTB + BL.GE + BL..BOP + BL..PD.4 + BL..CAL.3,  
  
    data = data,  
    family = "binomial"  
  )  
  
borthweight_model <-  
  lm(  
    Birthweight ~  
    Group *(Race_ethnicity + Public.Asstce + Use.Tob +  
            N.prev.preg + Live.PTB + BL.GE + BL..BOP +  
            BL..PD.4 + BL..CAL.3),  
  
    data = data  
  )
```

We consider less flexible regression methods in order to allow a higher chance at specifying at least one model for the doubly robust estimator. We wish to observe that a point estimate for ATE will be consistent with IPW2, but will have a smaller bootstrapped SE. My code for obtaining 500 bootstrap iterations is given in an appendix.

My code for calculating ATE using AIPW is given below. I am giving a step by step estimation of ATE for both outcomes variables, and will comment of the results one at a time.

```
# Step 1: get weights  
prop_scores = logistic_regression_a$fitted.values  
data_1 <- data_0 <- data  
  
data_1$Group = "T"  
data_0$Group = "C"  
  
preg_y_1 = predict(object = pregnancy_model, newdata = data_1, type = "response")
```

```

preg_y_0 = predict(object = pregnancy_model, newdata = data_0, type = "response")

bw_y_1 = predict(object = borthweight_model, newdata = data_1)
bw_y_0 = predict(object = borthweight_model, newdata = data_0)

### record actual outcomes and treatment flags
preg_y = ifelse(data$Preg.ended...37.wk == "Yes", 1, 0)
bw_y = data$Birthweight
trt_ind = ifelse(data$Group == "T", 1, 0)

### Pregnancy ATE
trt_ind * preg_y / prop_scores -
  ((trt_ind - prop_scores)/prop_scores) * preg_y_1 -> E_Y_1

(1 - trt_ind) * preg_y / (1 - prop_scores) -
  ((1 - trt_ind) - (1 - prop_scores))/
  (1 - prop_scores) * preg_y_0 -> E_Y_0

AIPW_preg_ate = mean(E_Y_1 - E_Y_0)

### Birthweight ATE
trt_ind * bw_y / prop_scores -
  ((trt_ind - prop_scores)/prop_scores) * bw_y_1 -> E_Y_1

(1 - trt_ind) * bw_y / (1 - prop_scores) -
  ((1 - trt_ind) - (1 - prop_scores))/
  (1 - prop_scores) * bw_y_0 -> E_Y_0

AIPW_bw_ate = mean(E_Y_1 - E_Y_0)

```

Pre-term pregnancy AIPW

- Pre-term pregnancy rate reduction estimate is -0.03, suggesting that the program was responsible for a -3% reduction in the rates of pre-term pregnancies as a result of program intervention.
- Bootstrapped standard error is 0.03
- A normal 95% confidence interval for the treatment effect is given by (-0.08, 0.03)

Birth weight AIPW

- Average birth weight of newborns increased by 73.81 as a result of program intervention.
- Bootstrapped standard error is 55.61
- A normal 95% confidence interval for the treatment effect is given by (-35.17, 182.8)

Pre-term pregnancies: comparison with previous results




Figure 2: Comparison of Estimation Method Variances for Pre-term pregnancy rate reduction

Birth weights: comparison with previous results

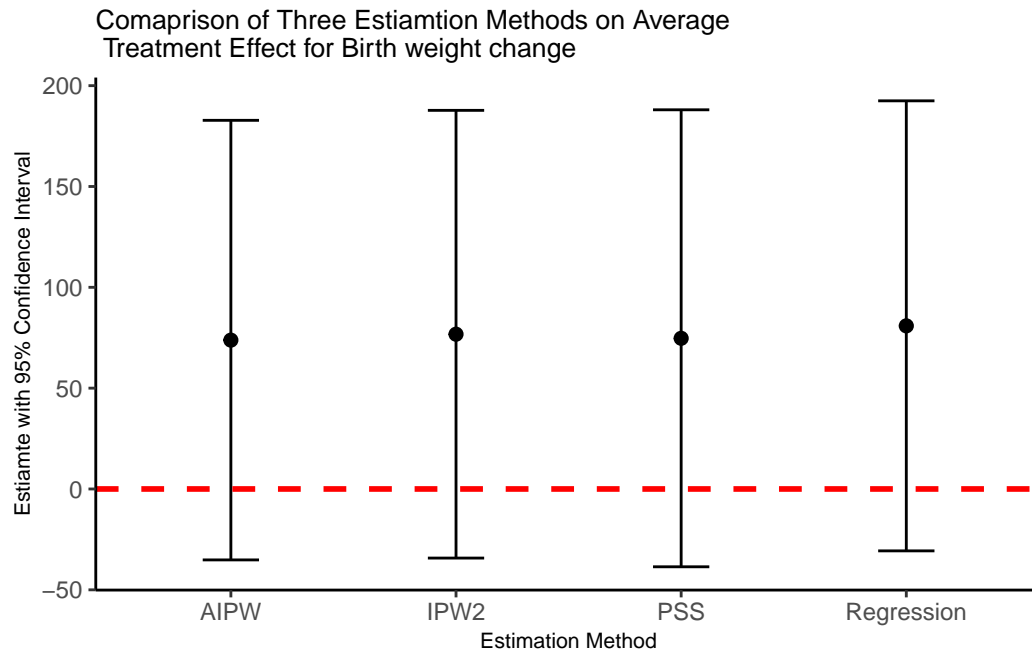


Figure 3: Comparison of Estiamtion Method Varinances for Average Birth Weight Increase

Appendix: bootsrap code

```
K = 500

res <-
  data.frame(
    i = 1:K,
    AIPW_preg_ate = NA,
    AIPW_bw_ate = NA
  )

set.seed(718297)
for(i in 1:K){

  print(i)
  iter_data = data[sample(1:nrow(data), replace = T), ]

  ##### train propensity score, pregnancy, and birth weights models
  iter_prop_score <-
    glm(I(iter_data$Group == "T") %>% as.numeric() ~
```

```

      Race_ethnicity + Public.Asstce +
      Use.Tob + Live.PTB +

      N.prev.preg + BL.GE + BL..BOP + BL..PD.4 + BL..CAL.3,

  data = iter_data,
  family = "binomial"
)

iter_pregnancy_model <-
  glm(
    `Preg.ended...37.wk` ~
      Group + Race_ethnicity + Public.Asstce + Use.Tob + N.prev.preg +
      Live.PTB + BL.GE + BL..BOP + BL..PD.4 + BL..CAL.3,

    data = iter_data,
    family = "binomial"
  )

iter_borthweight_model <-
  lm(
    Birthweight ~
      Group *(Race_ethnicity + Public.Asstce + Use.Tob +
        N.prev.preg + Live.PTB + BL.GE + BL..BOP +
        BL..PD.4 + BL..CAL.3),

    data = iter_data
  )

#### propensity scores, counterfactuals
prop_scores = iter_prop_score$fitted.values

iter_data_1 <- iter_data_0 <- iter_data

iter_data_1$Group = "T"
iter_data_0$Group = "C"

preg_y_1 = predict(object = iter_pregnancy_model, newdata = iter_data_1, type = "response")
preg_y_0 = predict(object = iter_pregnancy_model, newdata = iter_data_0, type = "response")

bw_y_1 = predict(object = iter_borthweight_model, newdata = iter_data_1)

```

```

bw_y_0 = predict(object = iter_borthweight_model, newdata = iter_data_0)

### record actual outcomes and treatment flags
preg_y = ifelse(iter_data$Preg.ended...37.wk == "Yes", 1, 0)
bw_y = iter_data$Birthweight
trt_ind = ifelse(iter_data$Group == "T", 1, 0)

### Pregnancy ATE
trt_ind * preg_y / prop_scores -
  ((trt_ind - prop_scores)/prop_scores) * preg_y_1 -> E_Y_1

(1 - trt_ind) * preg_y / (1 - prop_scores) -
  ((1 - trt_ind) - (1 - prop_scores))/
  (1 - prop_scores) * preg_y_0 -> E_Y_0

res$AIPW_preg_ate[i] = mean(E_Y_1 - E_Y_0)

### Birthweight ATE
trt_ind * bw_y / prop_scores -
  ((trt_ind - prop_scores)/prop_scores) * bw_y_1 -> E_Y_1

(1 - trt_ind) * bw_y / (1 - prop_scores) -
  ((1 - trt_ind) - (1 - prop_scores))/
  (1 - prop_scores) * bw_y_0 -> E_Y_0

res$AIPW_bw_ate[i] = mean(E_Y_1 - E_Y_0)

}

write.csv(res, "Augment IPW bootstrap results.csv")

```